



The two means over initialisations, superimposed and joined per mixture they differ because MDS is not linear; the rotation between the two fits is already removed

● after-embedding mean (centroid of the ten points) ▲ before-embedding mean (the ten matrices averaged, then embedded)


Structural
all layers · o_proj



Functional
all hidden states



Behavioral
R=16 per query



mixture g1/g2/g3 (%)

● 100/0/0	● 75/0/25	● 50/25/25	● 33/33/33	● 25/50/25	● 25/0/75	● 0/75/25	● 0/25/75
● 75/25/0	● 50/50/0	● 50/0/50	● 25/75/0	● 25/25/50	● 0/100/0	● 0/50/50	● 0/0/100